

RNA-seq Quality Assessment Assignment - Bi 623 (Summer 2021)



Data:

7_2E_fox_S6_L008



19_3F_fox_S14_L008

Part 1 – Read quality score distributions



7_2E_Fox Read 1 Graphs





7_2E_Fox Read 2 Graphs



19_3F_Fox Read 1 Graphs



19_3F_Fox Read 2 Graphs



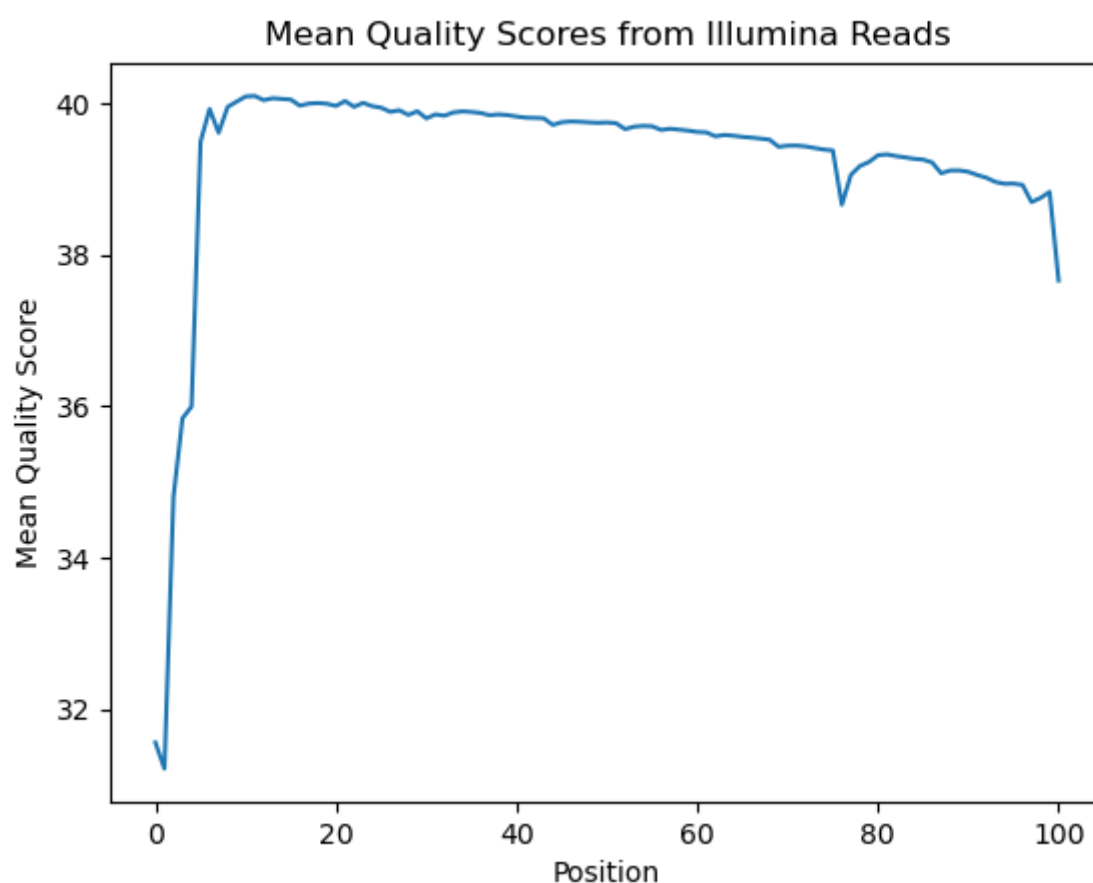
Answers

Are the per-base N plots consistent with the quality score plots?


The per-base N plots appear to be mostly consistent with the quality score plots. In all of the above plots we see a rise in the numebr of N's over the first few reads in the N content graphs. This is mirrored in the lower average quality score over these same reads in the quality score graphs.

Personal Quality Control Graphs


7_2E_Fox Read 1 Graphs




7_2E_Fox Read 2 Graphs



19_3F_Fox Read 1 Graph



19_3F_Fox Read 2 Graph



Answers

Describe how the FastQC quality score distribution plots compare to your own. If different, propose an explanation. Also, does the runtime differ? If so, why?

The quality plots are fairly similar to my own. Across all of the data both graphing methods show a clear decrease in quality scores in the early positions that quickly ramps after reading begins. In addition, both methods show a lower average data quality on 7_2E's 2nd read. Overall, I would say that these graphs mostly mirror each other, however the fastqc graphs present more information in a better way even if the core information is the same. Also, the runtime was significantly shorter for FastQC as compared to my own algorithm. If I had to guess, it would mostly be due to the inefficiency of looping through every line individually without indexing in my algorithm as compared to the FastQC method.

Part 2 – Adaptor trimming comparison

What proportion of reads (both forward and reverse) were trimmed?

```

7_2E read 1:
Total reads processed:          5,278,425
Reads with adapters:           173,473 (3.3%)
3.3% of total reads were trimmed

7_2E read 2:
Total reads processed:          5,278,425
Reads with adapters:           212,512 (4.0%)
4.0% of total reads were trimmed

19_3F read 1:
Total reads processed:          16,348,255
Reads with adapters:           546,623 (3.3%)
3.3% of total reads were trimmed

19_3F read 2:
Total reads processed:          16,348,255
Reads with adapters:           676,564 (4.1%)
4.1% of total reads were trimmed

```

Sanity check: Use your Unix skills to search for the adapter sequences in your datasets and confirm the expected sequence orientations. Report the commands you used, the reasoning behind them, and how you confirmed the adapter sequences.

A quick sanity check after googling the adapters that were used in the kit that this experiment with done with is shown below. The easiest way to see if these adapters show up in the data is a quick a grep. This will highlight where the sequence is found with in the file and will let us look at every line where it is found. We would expect that these adapters should appear after the insert is sequenced so we would not expect many instances of it at the start of the line and instead majority of them at the end. While this method doesn't allow us to check if the adapter is partially present in our reads it does give us a good indication that it is present in sufficient quantities in our reads.

Code examples:

This will visually show where the adapter is found in this data set.

```

zcat
/projects/bgmp/shared/2017_sequencing/demultiplexed/19_3F_fox_S14_L008_R1_0
01.fastq.gz | grep 'AGATCGGAAGAGCACACGTCTGAACTCCAGTCA'

```

This will tell us how much of the adapter was found at the start of a read.

```

zcat
/projects/bgmp/shared/2017_sequencing/demultiplexed/19_3F_fox_S14_L008_R1_0
01.fastq.gz | grep '^AGATCGGAAGAGCACACGTCTGAACTCCAGTCA' | wc -l

```

Plot the trimmed read length distributions for both R1 and R2 reads (on the same plot). You can produce 2 different plots for your 2 different RNA-seq samples. There are a number of ways you could possibly do this. One useful thing your plot should show, for example, is whether R1s are trimmed more extensively than R2s, or vice versa.

Comment on whether you expect R1s and R2s to be adapter-trimmed at different rates.

I would expect R2 to be trimmed more extensively than R1.

Part 3 – Alignment and strand-specificity

Mouse Genome Used

Genome name: `Mus_musculus.GRCm39.dna.primary_assembly.fa`

GTF name: `Mus_musculus.GRCm39.104.gtf`

Count reads that map to features using `htseq-count`. You should run `htseq-count` twice: once with `--stranded=yes` and again with `--stranded=no`. Use default parameters otherwise.

Demonstrate convincingly whether or not the data are from “strand-specific” RNA-Seq libraries. Include any commands/scripts used. Briefly describe your evidence, using quantitative statements (e.g. “I propose that these data are/are not strand-specific, because X% of the reads are y, as opposed to z.”).

Hint - recall ICA4 from Bi621.

To turn in your work for this assignment: Upload your Talapas batch script/code, FastQC plots, mapped/unmapped read counts, counts files generated from `htseq-count`, answers to questions, and any additional plots/code to github. You should create at most 2 files for submission (R markdown and the rendered pdf file) containing all these elements. The three parts of the assignment should be clearly labeled. Be sure to title and write a figure legend for each image/graph/table you present.